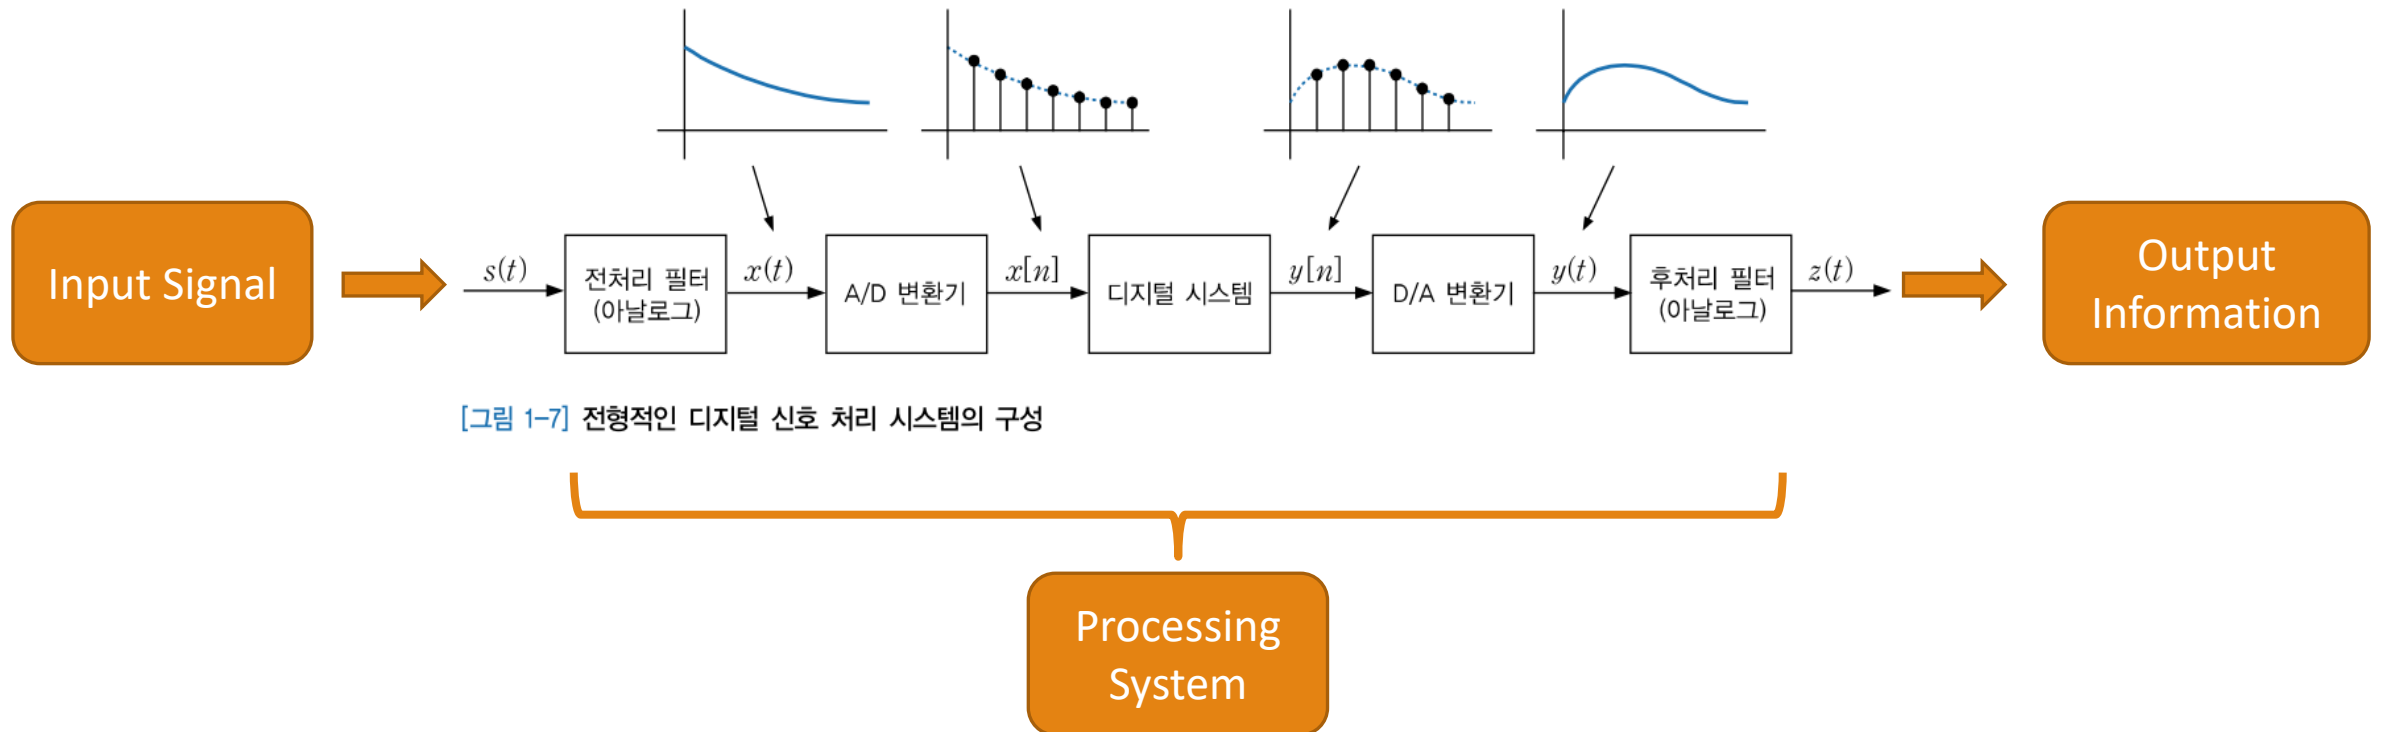


딥러닝을 활용한 디지털 영상 처리

Digital Image Processing via Deep Learning

Lecture 1 – Introduction to Image Processing and Classification

Thinking back to Signal Processing



How humans see the World



We observe the world via our eyes mostly.

Our eyes process the image as a signal and send it to our brain.

(Similar to pre-processing of a signal)


The brain receives the signal and determines what we are seeing.

A Cute Cat



Treat Image as a Signal

The computer treats image/images from a video as a matrix filled in with numbers representing grayscale intensity and RGB colour information



This image by Nikita is
licensed under [CC-BY 2.0](#)

```
[[105 112 108 111 104 99 106 99 96 103 112 119 104 97 93 87]  
[ 91 98 102 106 104 79 98 103 99 105 123 136 110 105 94 85]  
[ 76 85 90 105 128 105 87 96 95 99 115 112 106 103 99 85]  
[ 99 81 81 93 120 131 127 100 95 98 102 99 96 93 101 94]  
[106 91 61 64 69 91 88 85 101 107 109 98 75 84 96 95]  
[114 108 85 55 55 69 64 54 64 87 112 129 98 74 84 91]  
[133 137 147 103 65 81 80 65 52 54 74 84 102 93 85 82]  
[128 137 144 140 109 95 86 70 62 65 63 63 60 73 86 101]  
[125 133 148 137 119 121 117 94 65 79 80 65 54 64 72 98]  
[127 125 131 147 133 127 126 131 111 96 89 75 61 64 72 84]  
[115 114 109 123 150 148 131 118 113 109 100 92 74 65 72 78]  
[ 89 93 90 97 108 147 131 118 113 114 113 109 106 95 77 80]  
[ 63 77 86 81 77 79 102 123 117 115 117 125 125 130 115 87]  
[ 62 65 82 89 78 71 80 101 124 126 119 101 107 114 131 119]  
[ 63 65 75 88 89 71 62 81 120 138 135 105 81 98 110 118]  
[ 87 65 71 87 106 95 69 45 76 130 126 107 92 94 105 112]  
[118 97 82 86 117 123 116 66 41 51 95 93 89 95 102 107]  
[164 146 112 80 82 120 124 104 76 48 45 66 88 101 102 109]  
[157 170 157 120 93 86 114 132 112 97 69 55 70 82 99 94]  
[130 128 134 161 139 100 109 118 121 134 114 87 65 53 69 86]  
[128 112 96 117 150 144 120 115 104 107 102 93 87 81 72 79]  
[123 107 96 86 83 112 153 149 122 109 104 75 80 107 112 99]  
[122 121 102 80 82 86 94 117 145 148 153 102 58 78 92 107]  
[122 164 148 103 71 56 78 83 93 103 119 139 102 61 69 84]]
```


What the computer sees

An image is just a big grid of numbers between [0, 255]:

e.g. 800 x 600 x 3
(3 channels RGB)

Image Processing Techniques

Histogram Equalization



Edge Detection



Feature point extraction




Image Classification

Knowing what is in the image / what the image indicates is a major challenge in computer vision

Binary Classification



Multiclass Classification



Multilabel Classification






Image Classification

```
def classify_image(image):  
    # Some magic here?  
    return class_label
```

Input Image to be classified



Output label

Cat

Image Classification

How can we set up the magic box?

Attempt 1: Find edges of the image and use feature extractor to identify shapes

Drawback: Extremely complex algorithms and set of rules are different for all features.




Image Classification

Attempt 2: Data Driven Approach

10 classes

50,000 training images

10,000 testing images



Test images and nearest neighbors



Image Classification: K Nearest Neighbor

K Nearest Neighbor (KNN):

- For training, simply memorise all training data.
- For testing, for each test image, find the closest train image and predict the label of the nearest image




Image Classification: K Nearest Neighbor


L1 Distance: $d_{L1}(I_1, I_2) = \sum_i |I_1^i - I_2^i|$

L2 Distance: $d_{L2}(I_1, I_2) = \sqrt{\sum_i (I_1^i - I_2^i)^2}$, where I_n is a set of images and i indicates for a class.


test image					training image					pixel-wise absolute value differences				
56	32	10	18		10	20	24	17		46	12	14	1	
90	23	128	133		8	10	89	100		82	13	39	33	
24	26	178	200	-	12	16	178	170	=	12	10	0	30	
2	0	255	220		4	32	233	112		2	32	22	108	add → 456

Image Classification: K Nearest Neighbor


KNN takes the majority votes from K nearest points



$K = 1$



$K = 3$



$K = 5$

Try yourself: <http://vision.stanford.edu/teaching/cs231n-demos/knn/>

Setting Hyperparameters

What is the best K?

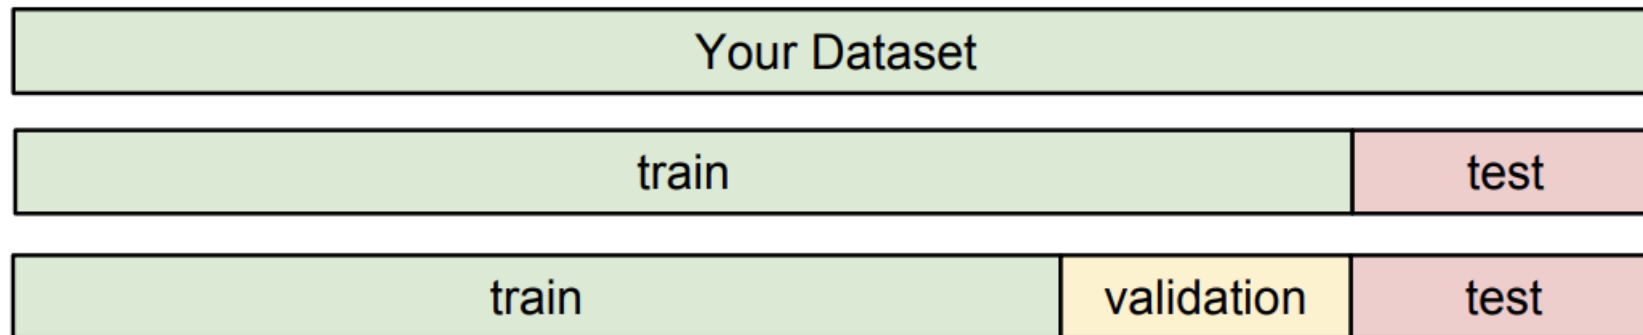
What is the best distance metric to use?

These are hyperparameters: choice about the algorithm that we set rather than learn.

The hyperparameters are very problem-dependent. Must try them all out and see what works the best.

Given a Dataset, we usually leave a portion out for testing.

In addition, we can leave a validation set for selecting appropriate hyperparameters.



Setting Hyperparameters

Cross-Validation: Split the data into folds, try each fold as validation, and average the results.

It is useful for small datasets, but not frequently used in Deep Learning

fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test

Image Classification

The idea of KNN is good but KNN is actually never used in image classification.

- Very slow at test time: We want an algorithm that may be slow during training but fast during testing. KNN is fast during training since it simply memorises data but slow during testing due to distance computation.
- The Distance metrics on pixels are not informative



(all 3 images have same L2 distance to the one on the left)

Image Classification: Linear Classification

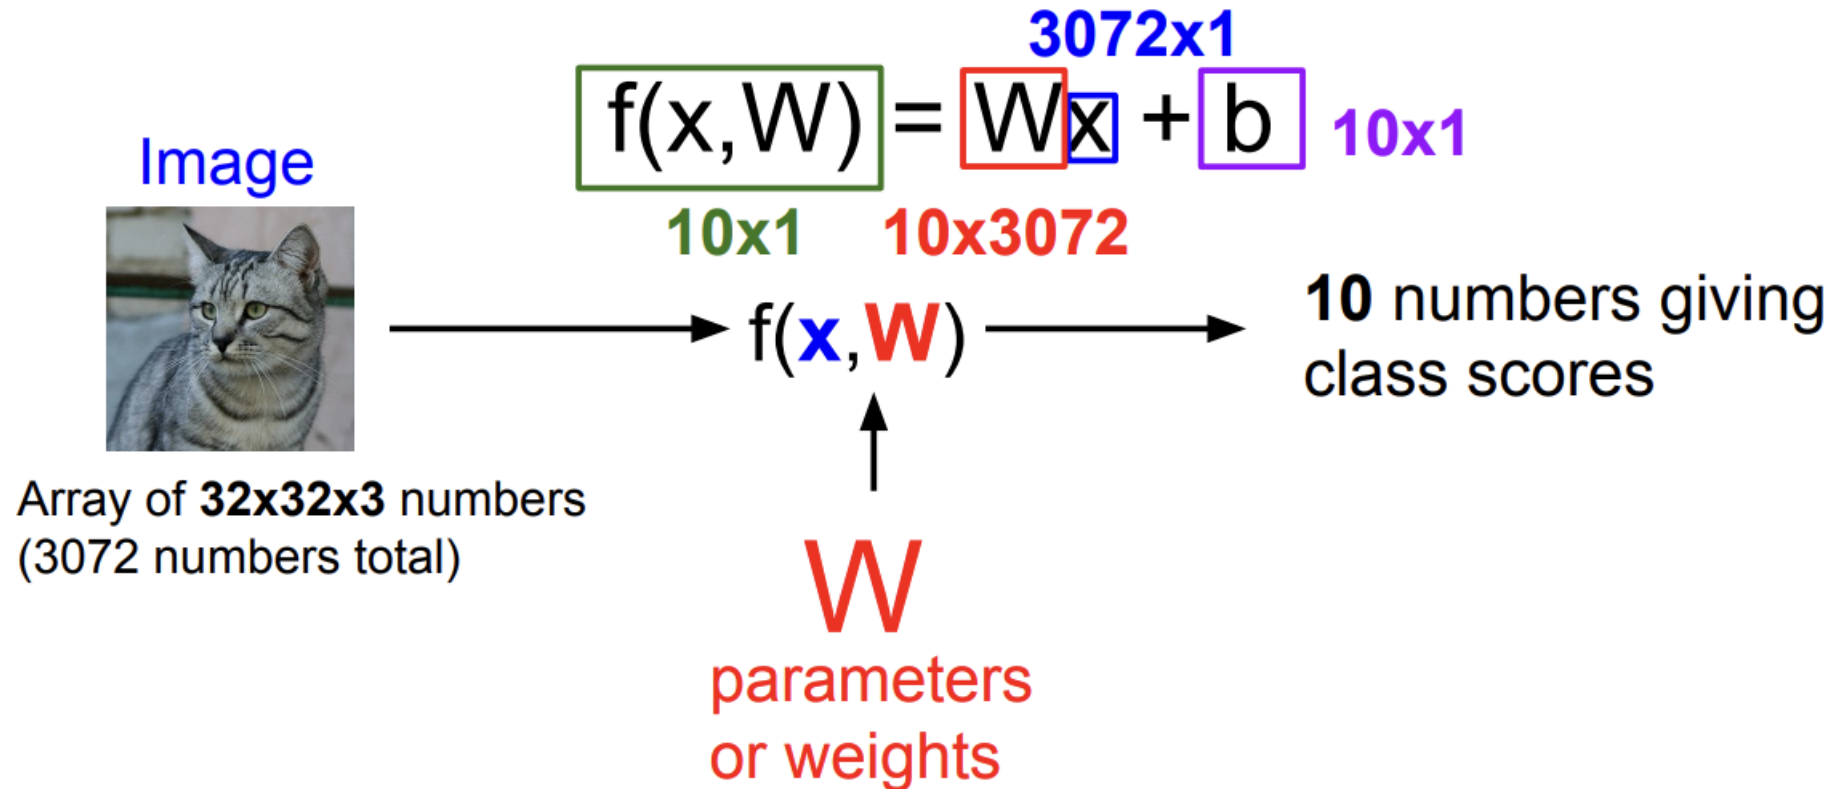


Image Classification: Linear Classification

Example with an image with 4 pixels and 3 classes (cat / dog / ship)

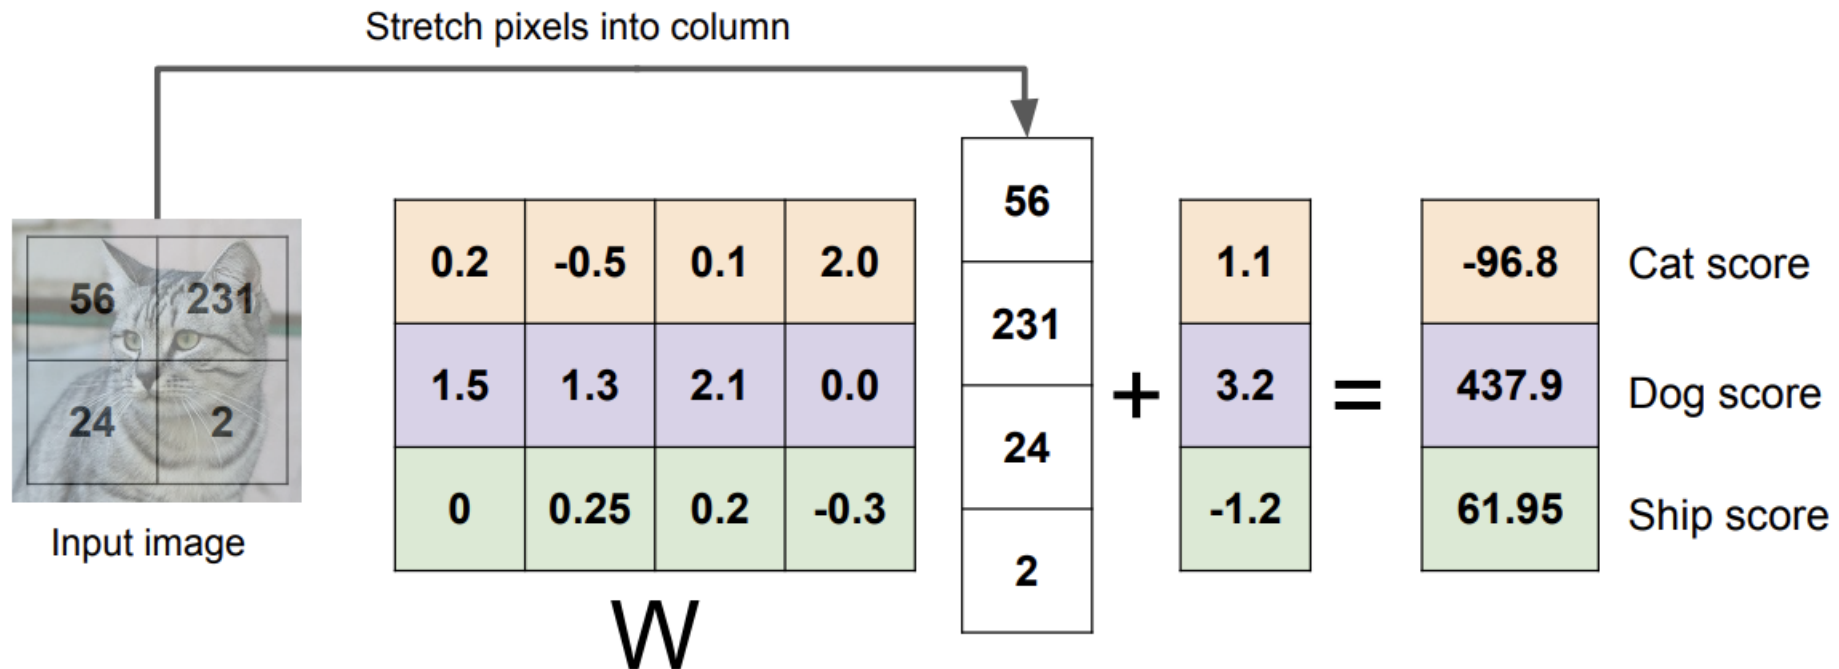
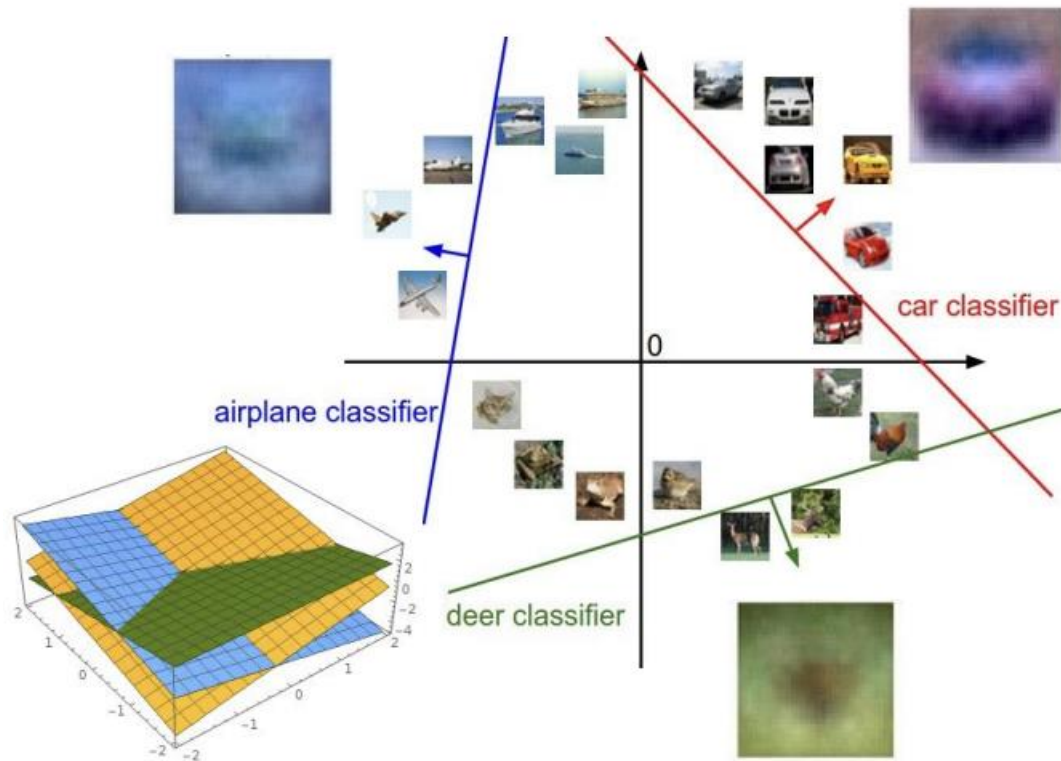


Image Classification: Linear Classification



$$f(x, W) = Wx + b$$



Array of **32x32x3** numbers
(3072 numbers total)

Image Classification: Linear Classification

Example
Scores for 3
images



airplane	-3.45	-0.51	3.42
automobile	-8.87	6.04	4.64
bird	0.09	5.31	2.65
cat	2.9	-4.22	5.1
deer	4.48	-4.19	2.64
dog	8.02	3.58	5.55
frog	3.78	4.49	-4.34
horse	1.06	-4.37	-1.5
ship	-0.36	-2.09	-4.79
truck	-0.72	-2.93	6.14

Linear Regression: Polynomial Curve Fitting


Green Curve: function used to generate data $\sin(2\pi x)$

Blue Points: Dataset generated via the green curve with some noise

Goal: Predict a function that fits the blue points well and represents the green curve as closely as possible, assuming we do not know the green curve.

A Polynomial function is common to start with:

$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$




Linear Regression: Polynomial Curve Fitting

We measure and error, the difference between the data value (blue point) and the prediction, y

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$

Our objective is to find the weights, \mathbf{w} , that minimises the error.







Linear Regression: Polynomial Curve Fitting

The results differ by the value of M chosen, which is the order of the polynomial

Which one seems to fit best?

$M = 0$ & $M = 1$ is known as **underfitting**. It hardly represents the pattern of data.

$M = 9$ is known as **overfitting**. The error is 0 but cannot predict the value outside the training set.




Linear Regression: Polynomial Curve Fitting

Overfitting occurs when the order of the polynomial is high compared to the number of training data.



	$M = 0$	$M = 1$	$M = 3$	$M = 9$
w_0^*	0.19	0.82	0.31	0.35
w_1^*		-1.27	7.99	232.37
w_2^*			-25.43	-5321.83
w_3^*			17.37	48568.31
w_4^*				-231639.30
w_5^*				640042.26
w_6^*				-1061800.52
w_7^*				1042400.18
w_8^*				-557682.99
w_9^*				125201.43

For higher order polynomials, the weights seem to have large absolute values.



Linear Regression: Polynomial Curve Fitting

Solution to Overfitting 1: Increase the number of training dataset



Linear Regression: Polynomial Curve Fitting

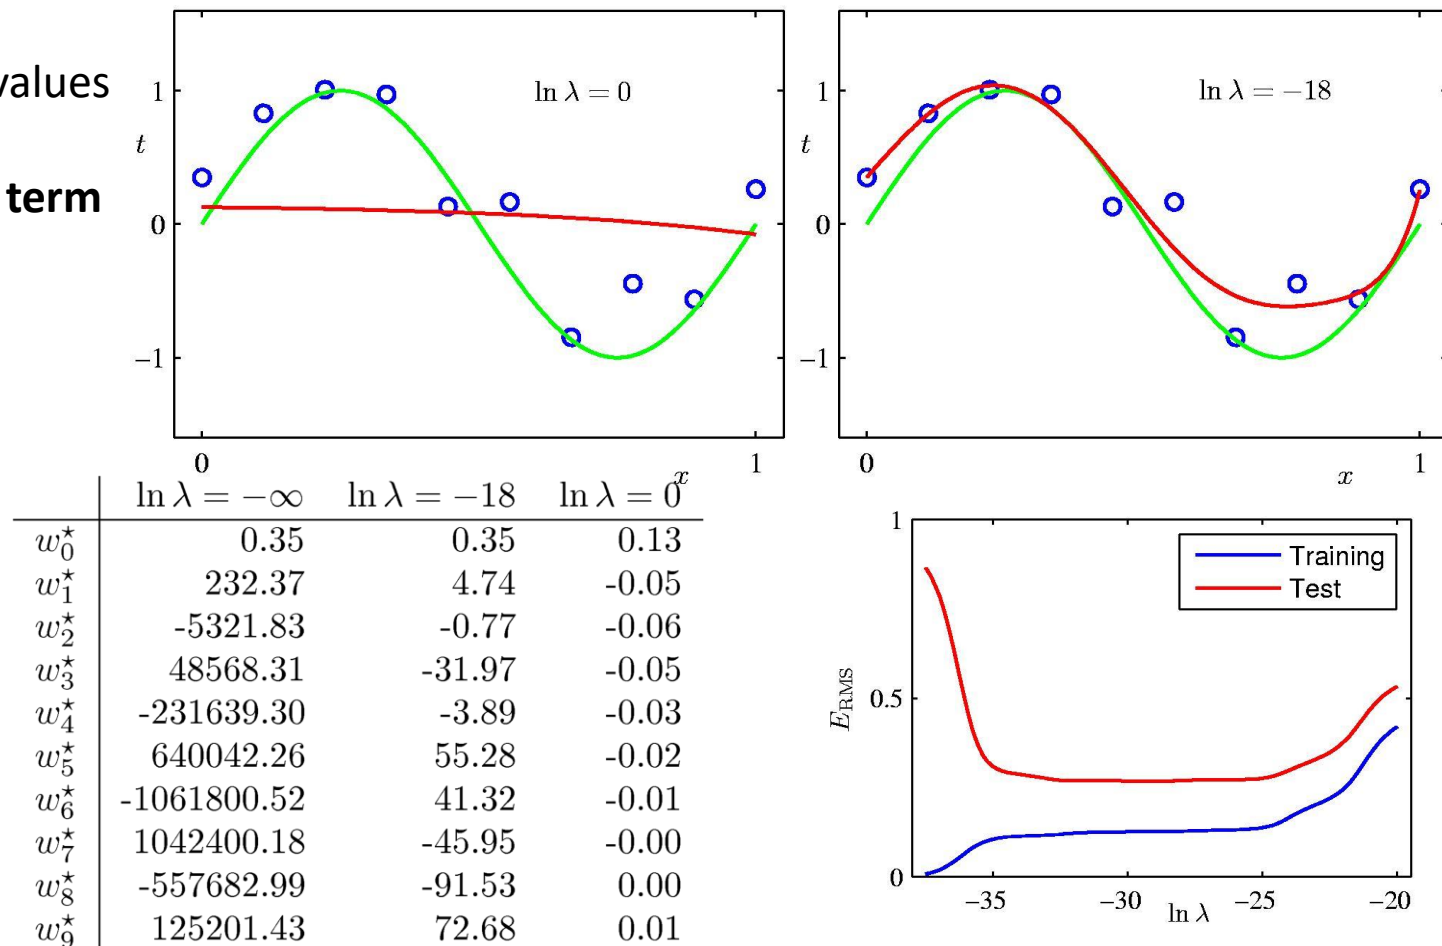
Solution to Overfitting 2: Penalize for large coefficient values

Rewrite the error function including the **regularization term**

This is called the **ridge regression**

$$\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

$\lambda = 0$ ($\ln(\lambda) = -\infty$) corresponds to no regularization.



Linear Regression: Polynomial Curve Fitting

How do we find the optimal weights?

$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$

$$\mathbf{w} = \operatorname{argmin}_{\mathbf{w}} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 = \operatorname{argmin}_{\mathbf{w}} (\mathbf{x}\mathbf{w} - \mathbf{t})^T (\mathbf{x}\mathbf{w} - \mathbf{t}) \longrightarrow \boxed{\mathbf{w} = (\mathbf{x}^T \mathbf{x})^{-1} \mathbf{x}^T \mathbf{y}}$$

The complexity of this computation is $O(n^3)$

As the size of the dataset increase, weights become extremely difficult to find.
We must seek more convenient methods to find the weights

Towards Deep Learning

Can Linear Classifier classify all the images? How about non-linear relationships?

How can we find appropriate weights not by heavy computation?

Deep Learning methods remedy (partially in some sense) these shortcomings.

Next Lecture

Perceptron and Introduction to Deep Learning

